

---

# CS481/CS583: Bioinformatics Algorithms

---

Can Alkan

EA509

`calkan@cs.bilkent.edu.tr`

<http://www.cs.bilkent.edu.tr/~calkan/teaching/cs481/>

---

# **SIMILARITY SEARCH**

---

# Heuristic Similarity Searches

- Genomes are huge: Smith-Waterman quadratic alignment algorithms are too slow
- Alignment of two sequences usually has short identical or highly similar fragments
- Many heuristic methods (i.e., FASTA) are based on the same idea of *filtration*
  - Find short exact matches, and use them as seeds for potential match extension
  - “Filter” out positions with no extendable matches


# Dot Matrices

- Dot matrices show similarities between two sequences
- FASTA makes an implicit dot matrix from short exact matches, and tries to find long diagonals (allowing for some mismatches)

|   | G | A | T | T | C | G | C | T | T | A | G | T |
|---|---|---|---|---|---|---|---|---|---|---|---|---|
| C |   |   |   |   | * |   | * |   |   |   |   |   |
| T |   | * | * |   |   |   |   | * | * |   |   | * |
| G | * |   |   |   |   | * |   |   |   |   |   | * |
| A |   | * |   |   |   |   |   |   |   | * |   |   |
| T |   | * | * |   |   |   | * | * |   | * |   | * |
| T |   | * | * |   |   |   | * | * |   | * |   | * |
| C |   |   |   |   | * |   | * |   |   |   |   |   |
| C |   |   |   |   | * |   | * |   |   |   |   |   |
| T |   | * | * |   |   |   | * | * |   | * |   | * |
| T |   | * | * |   |   |   | * | * |   | * |   | * |
| A |   | * |   |   |   |   |   |   |   | * |   |   |
| G | * |   |   |   |   | * |   |   |   |   | * |   |
| T |   | * | * |   |   |   | * | * |   | * |   | * |
| C |   |   |   |   | * |   | * |   |   |   |   |   |
| A |   | * |   |   |   |   |   |   |   | * |   |   |
| G | * |   |   |   |   | * |   |   |   |   | * |   |


# Dot Matrices (cont'd)

- Identify diagonals above a threshold length
- Diagonals in the dot matrix indicate exact substring matching



# Diagonals in Dot Matrices

- Extend diagonals and try to link them together, allowing for minimal mismatches/indels
- Linking diagonals reveals approximate matches over longer substrings



# Approximate Pattern Matching Problem

- Goal: *Find all approximate occurrences of a pattern in a text*
- Input: A pattern  $\mathbf{p} = p_1 \dots p_n$ , text  $\mathbf{t} = t_1 \dots t_m$ , and  $k$ , the maximum number of mismatches
- Output: All positions  $1 \leq i \leq (m - n + 1)$  such that  $t_i \dots t_{i+n-1}$  and  $p_1 \dots p_n$  have at most  $k$  mismatches (i.e., Hamming distance between  $t_i \dots t_{i+n-1}$  and  $\mathbf{p} \leq k$ )

# Approximate Pattern Matching: A Brute-Force Algorithm

## ApproximatePatternMatching(**p**, **t**, *k*)

```
1  n  $\square$  length of pattern p
2  m  $\square$  length of text t
3  for i  $\square$  1 to m - n + 1
4      dist  $\square$  0
5      for j  $\square$  1 to n
6          if  $t_{i+j-1} \neq p_j$ 
7              dist  $\square$  dist + 1
8      if dist  $\leq k$ 
9          output i
```



## Approximate Pattern Matching: Running Time

- That algorithm runs in  $O(nm)$ .
- We can generalize the “Approximate Pattern Matching Problem” into a “Query Matching Problem”:
  - We want to match substrings in a query to substrings in a text with at most  **$k$**  mismatches
  - **Motivation:** we want to see similarities to some gene, but we may not know which parts of the gene to look for

# Query Matching Problem

- Goal: Find all substrings of the query that approximately match the text
- Input: Query  $\mathbf{q} = q_1 \dots q_w$ ,  
text  $\mathbf{t} = t_1 \dots t_m$ ,  
 $n$  (length of matching substrings),  
 $k$  (maximum number of mismatches)
- Output: All pairs of positions  $(i, j)$  such that the  
 $n$ -letter substring of  $\mathbf{q}$  starting at  $i$   
approximately matches the  
 $n$ -letter substring of  $\mathbf{t}$  starting at  $j$ ,  
with at most  $k$  mismatches

---

# Query Matching: Main Idea

- Approximately matching strings share some perfectly matching substrings.
  - Instead of searching for approximately matching strings (difficult) search for perfectly matching substrings (easy).
-

# Filtration in Query Matching


- We want all  $n$ -matches between a query and a text with up to  $k$  mismatches
- “Filter” out positions we know do not match between text and query
- **Potential match detection**: find all matches of  $\ell$ -tuples in query and text for some small  $\ell$
- **Potential match verification**: Verify each potential match by extending it to the left and right, until  $(k + 1)$  mismatches are found

# Filtration: Match Detection

- If  $x_1 \dots x_n$  and  $y_1 \dots y_n$  match with at most  $k$  mismatches, they must share an  $\ell$ -tuple that is perfectly matched, with  $\ell = \lfloor n/(k+1) \rfloor$
- Break string of length  $n$  into  $k+1$  parts, each each of length  $\lfloor n/(k+1) \rfloor$ 
  - $k$  mismatches can affect at most  $k$  of these  $k+1$  parts
  - At least one of these  $k+1$  parts is perfectly matched

# Filtration: Match Detection (cont'd)

- Suppose  $k = 3$ . We would then have  $l = n/(k+1) = n/4$ :




- There are at most  $k$  mismatches in  $n$ , so at the very least there must be one out of the  $k+1$   $l$ -tuples without a mismatch

What is this based on?

# Filtration: Match Verification

- For each  $\ell$ -match we find, try to extend the match further to see if it is substantial



# Filtration: Example

|                      | $k = 0$ | $k = 1$ | $k = 2$ | $k = 3$ | $k = 4$ | $k = 5$ |
|----------------------|---------|---------|---------|---------|---------|---------|
| $\ell$ -tuple length | $n$     | $n/2$   | $n/3$   | $n/4$   | $n/5$   | $n/6$   |

Shorter perfect matches required

Performance decreases



---

Lipman & Pearson, 1985

**FASTP**

---

# FASTP

- Three phase algorithm
  1. Find short good matches using k-mers
    1.  $k=1, k=2$
  2. Find start and end positions for good matches
  3. Use DP to align good matches

# FASTP: Phase 1 (1)

```
position  1 2 3 4 5 6 7 8 9 10 11
protein 1 n c s p t a . . . . .
protein 2 . . . . . a c s p r k
```


| amino acid | position in |           | offset       |
|------------|-------------|-----------|--------------|
|            | protein 1   | protein 2 | pos 1 - pos2 |
| a          | 6           | 6         | 0            |
| c          | 2           | 7         | -5           |
| k          | -           | 11        |              |
| n          | 1           | -         |              |
| p          | 4           | 9         | -5           |
| r          | -           | 10        |              |
| s          | 3           | 8         | -5           |
| t          | 5           | -         |              |

Note the common offset for the 3 amino acids c,s and p  
A possible alignment can be quickly found :

```
protein 1 n c s p t a
           | | |
protein 2 a c s p r k
```


# FASTP: Phase 1 (2)

- Similar to dot plot
- Offsets range from 1-m to n-1
- Each offset is scored as
  - $\# \text{ matches} - \# \text{ mismatches}$
- Diagonals (offsets) with large score show local similarities



# FASTP: Phase 2

- 5 best diagonal runs are found
- Rescore these 5 regions using PAM250.
  - Initial score
- Indels are not considered yet



---

# FASTP: Phase 3

- Sort the aligned regions in descending score
  - Optimize these alignments using Needleman-Wunsch
  - Report the results
-

---


Pearson 1995

# **FASTA – IMPROVEMENT OVER FASTP**

---

# FASTA (1)

- Phase 2: Choose 10 best diagonal runs instead of 5






# FASTA (2)

## ■ Phase 2.5

- Eliminate diagonals that score less than some given threshold.
- Combine matches to find longer matches. It incurs join penalty similar to gap penalty



---

# FASTA Variations

- TFASTAX and TFASTAY: query protein against a DNA library in all reading frames
  - FASTAX, FASTAY: DNA query in all reading frames against protein database
-

---

# BLAST

---

# Local alignment is too slow...

- Quadratic local alignment is too slow while looking for similarities between long strings (e.g. the entire GenBank database)

$$s_{i,j} = \max \begin{cases} 0 \\ s_{i-1,j} + \delta(v_i, -) \\ s_{i,j-1} + \delta(-, w_j) \\ s_{i-1,j-1} + \delta(v_i, w_j) \end{cases}$$

# Local alignment is too slow...

- Quadratic local alignment is too slow while looking for similarities between long strings (e.g. the entire GenBank database)
- Guaranteed to find the optimal local alignment
- Sets the standard for sensitivity

$$s_{i,j} = \max \begin{cases} 0 \\ s_{i-1,j} + \delta(v_i, -) \\ s_{i,j-1} + \delta(-, w_j) \\ s_{i-1,j-1} + \delta(v_i, w_j) \end{cases}$$

# Local alignment is too slow...

- Quadratic local alignment is too slow while looking for similarities between long strings (e.g. the entire GenBank database)
- **Basic Local Alignment Search Tool**
  - Altschul, S., Gish, W., Miller, W., Myers, E. & Lipman, D.J.  
Journal of Mol. Biol., 1990
- Search sequence databases for local alignments to a query

$$s_{i,j} = \max \begin{cases} 0 \\ s_{i-1,j} + \delta(v_i, -) \\ s_{i,j-1} + \delta(-, w_j) \\ s_{i-1,j-1} + \delta(v_i, w_j) \end{cases}$$

# BLAST

- Great improvement in speed, with a modest decrease in sensitivity
- Minimizes search space instead of exploring entire search space between two sequences
- Finds short exact matches (“seeds”), only explores locally around these “hits”
  - “Seed-and-extend”

---

# What Similarity Reveals

- BLASTing a new gene
    - Evolutionary relationship
    - Similarity between protein function
  - BLASTing a genome
    - Potential genes
-



# BLAST algorithm

- **Keyword search** of all words of length  $w$  from the query of length  $n$  in database of length  $m$  with score above threshold
  - $w = 11$  for DNA queries,  $w = 3$  for proteins
  - For each  $k$ -mer  $w$  find all  $k$ -mer that aligns with score at least cutoff  $T$
- **Local alignment extension** for each found keyword
  - Extend result until longest match above threshold is achieved
- Running time  $O(nm)$

# BLAST algorithm (cont'd)

keyword

Query: KRHRKVLRDNIQGITKPAIRRLARRGGVKRISGLIYEETRGVLKIFLENVIRD

neighborhood  
score  
threshold  
(T = 13)

|     |    |
|-----|----|
| GVK | 18 |
| GAK | 16 |
| GIK | 16 |
| GGK | 14 |
| GLK | 13 |
| GNK | 12 |
| GRK | 11 |
| GEK | 11 |
| GDK | 11 |

Neighborhood  
words

extension

Query: 22 VLRDNIQGITKPAIRRLARRGGVKRISGLIYEETRGVLK 60

+++DN +G + IR L G+K I+ L+ E+ RG++K

Sbjct: 226 IIKDNGRGFSGKQIRNLNYGIGLKVIADLV-EKHRGIIK 263

High-scoring Pair (HSP)

---

# Original BLAST

- **Dictionary**

- All words of length  $w$

- **Alignment**


- Ungapped extensions until score falls below some statistical threshold

- **Output**

- All local alignments with score  $>$  threshold
-


# Original BLAST: Example

- $w = 4$
- Exact keyword match of GGTC
- Extend diagonals with mismatches until score is under 50%
- Output result  
GTAAGGTCC  
GTTAGGTCC



# Gapped BLAST : Example

- Original BLAST exact keyword search, THEN:
  - Extend with gaps around ends of exact match **until score < threshold**
  - Output result
- GTAAGGTCCAGT**  
**GTTAGGTC-AGT**



---

# Incarnations of BLAST

- blastn: Nucleotide-nucleotide
  - blastp: Protein-protein
  - blastx: Translated query vs. protein database
  - tblastn: Protein query vs. translated database
  - tblastx: Translated query vs. translated database (6 frames each)
-

---

# Incarnations of BLAST (cont'd)

- PSI-BLAST
    - Find members of a protein family or build a custom position-specific score matrix
  - Megablast:
    - Search longer sequences with fewer differences
  - WU-BLAST: (Wash U BLAST)
    - Optimized, added features
-

---

# ASSESSING SEQUENCE SIMILARITY

---



# Assessing sequence similarity

- Need to know how strong an alignment can be expected from chance alone
- “Chance” relates to comparison of sequences that are generated randomly based upon a certain sequence model
- Sequence models may take into account:
  - ❑ G+C content
  - ❑ Poly-A tails
  - ❑ “Junk” DNA
  - ❑ Codon bias
  - ❑ Etc.

# BLAST: Segment Score

- BLAST uses scoring matrices ( $\delta$ ) to improve on efficiency of match detection
  - Some proteins may have very different amino acid sequences, but are still similar
- For any two  $\ell$ -mers  $x_1 \dots x_\ell$  and  $y_1 \dots y_\ell$ :
  - Segment pair: pair of  $\ell$ -mers, one from each sequence
  - Segment score:  $\sum_{i=1}^{\ell} \delta(x_i, y_i)$

# BLAST: Locally Maximal Segment Pairs

- A segment pair is maximal if it has the best score over all segment pairs
- A segment pair is locally maximal if its score can't be improved by extending or shortening
- Statistically significant *locally maximal* segment pairs are of biological interest
- BLAST finds all locally maximal segment pairs with scores above some threshold
  - A significantly high threshold will filter out some statistically insignificant matches

# BLAST: Statistics

- Threshold: Altschul-Dembo-Karlin statistics
  - Identifies smallest segment score that is unlikely to happen by chance
- # matches above  $\theta$  has mean  $E(\theta) = Kmne^{-\lambda\theta}$ ;  $K$  is a constant,  $m$  and  $n$  are the lengths of the two compared sequences
  - Parameter  $\lambda$  is positive root of:  
$$\sum_{x,y \in A} (p_x p_y e^{\delta(x,y)}) = 1$$
, where  $p_x$  and  $p_y$  are frequencies of amino acids  $x$  and  $y$ , and  $A$  is the twenty letter amino acid alphabet

# P-values

- The probability of finding  $b$  high-scoring segment pairs (HSPs) with a score  $\geq S$  is given by:
  - $(e^{-E}E^b)/b!$
- For  $b = 0$ , that chance is:
  - $e^{-E}$
- Thus the probability of finding at least one HSP with a score  $\geq S$  is:
  - $P = 1 - e^{-E}$

# Sample BLAST output

- Blast of human beta globin protein against zebra fish**

Sequences producing significant alignments:

|  | Score  | E     |
|--|--------|-------|
|  | (bits) | Value |

|  |               |     |       |
|--|---------------|-----|-------|
| gi 18858329 ref NP_571095.1  ba1 globin [Danio rerio]                | >gi 147757... | 171 | 3e-44 |
| gi 18858331 ref NP_571096.1  ba2 globin; SI:dZ118J2.3 [Danio rer...] |               | 170 | 7e-44 |
| gi 37606100 emb CAE48992.1  SI:bY187G17.6 (novel beta globin) [D...] |               | 170 | 7e-44 |
| gi 31419195 gb AAH53176.1  Ba1 protein [Danio rerio]                 |               | 168 | 3e-43 |

## ALIGNMENTS

>gi|18858329|ref|NP\_571095.1| ba1 globin [Danio rerio]  
Length = 148

Score = 171 bits (434), Expect = 3e-44

Identities = 76/148 (51%), Positives = 106/148 (71%), Gaps = 1/148 (0%)

Query: 1 MVHLTPEEKSAVTALWGKVVNDEVGGEALGRLLVVYPWTQRFFESFGDLSTPDAVMGNPK 60

MV T E++A+ LWGK+N+DE+G +AL R L+VYPWTQR+F +FG+LS+P A+MGNPK

Sbjct: 1 MVEWTD AERTAILGLWGKLNIDEIGPQALSRC L I VYPWTQRYFATFGNLSSPAAIMGNPK 60

Query: 61 VKAHGKKVLGAFSDGLAHL DNLKGT FATLSELHCDKLHVDPENFRLLGNVLVCVLAHHFG 120

V AHG+ V+G + ++DN+K T+A LS +H +KLHVDP+NFRL L + + A FG

Sbjct: 61 VAAHGRTVMGGLERAIKNMDNVKNTYAALSV MHSEKLHVDPDNFRLLADCITVCAAMKFG 120

Query: 121 KE-FTPPVQAAYQKVVAGVANALAHKYH 147

+ F VQ A+QK +A V +AL +YH

Sbjct: 121 QAGFNADVQEA WQKFLAVVVSALCRQYH 148

# Sample BLAST output (cont'd)

- Blast of human beta globin DNA against human DNA**

Sequences producing significant alignments:

|  | Score<br>(bits) | E<br>Value |
|--|-----------------|------------|
|--|-----------------|------------|

|   |     |       |
|---|-----|-------|
| gi 19849266 gb AF487523.1  Homo sapiens gamma A hemoglobin (HBG1... | 289 | 1e-75 |
| gi 183868 gb M11427.1 HUMHBG3E Human gamma-globin mRNA, 3' end      | 289 | 1e-75 |
| gi 44887617 gb AY534688.1  Homo sapiens A-gamma globin (HBG1) ge... | 280 | 1e-72 |
| gi 31726 emb V00512.1 HSGGL1 Human messenger RNA for gamma-globin   | 260 | 1e-66 |
| gi 38683401 ref NR_001589.1  Homo sapiens hemoglobin, beta pseud... | 151 | 7e-34 |
| gi 18462073 gb AF339400.1  Homo sapiens haplotype PB26 beta-glob... | 149 | 3e-33 |

## ALIGNMENTS

>gi|28380636|ref|NG\_000007.3| Homo sapiens beta globin region (HBB@) on chromosome 11  
Length = 81706

Score = 149 bits (75), Expect = 3e-33

Identities = 183/219 (83%)

Strand = Plus / Plus

Query: 267     ttgggagatgccacaaagcacctggatgatctcaagggcacctttgcccagctgagtga 326

|| ||| | ||     | || |     ||||| ||||| ||||| ||||| |||||

Sbjct: 54409    ttcgaaaagctgttatgctcacggatgacctcaaaggcacctttgctacactgagtga 54468

Query: 327     ctgcactgtgacaagctgcatgtggatcctgagaacttc 365

||||||| ||||| ||||| ||||| ||||| ||||| |||||

Sbjct: 54469    ctgcactgtaacaagctgcacgtggaccctgagaacttc 54507

---

# Timeline

- 1970: Needleman-Wunsch global alignment algorithm
  - 1981: Smith-Waterman local alignment algorithm
  - 1985: FASTA
  - 1990: BLAST (basic local alignment search tool)
  - 2000s: BLAST has become too slow in “genome vs. genome” comparisons - new faster algorithms evolve!
    - Pattern Hunter
    - BLAT
-



---

Slides by Rayan Chikhi  
Institut Pasteur, CNRS

Original work by Heng Li

**Minimap2**

---

# Two articles

*Bioinformatics*, 32(14), 2016, 2103–2110

doi: 10.1093/bioinformatics/btw152

Advance Access Publication Date: 19 March 2016

Original Paper

OXFORD

---

Sequence analysis

## **Minimap and miniasm: fast mapping and de novo assembly for noisy long sequences**

**Heng Li**

Medical Population Genetics, Broad Institute, Cambridge, MA 02142, USA

Associate Editor: Inanc Birol

Received on December 6, 2015; revised on March 14, 2016; accepted on March 14, 2016

---

# Two articles

*Bioinformatics*, 34(18), 2018, 3094–3100

doi: 10.1093/bioinformatics/bty191

Advance Access Publication Date: 10 May 2018

Original Paper

OXFORD

---

Sequence analysis

## **Minimap2: pairwise alignment for nucleotide sequences**

**Heng Li\***

Department of Medical Population Genetics Program, Broad Institute, Cambridge, MA 02142, USA

\*To whom correspondence should be addressed.

Associate Editor: Inanc Birol

Received on January 1, 2018; revised on March 16, 2018; editorial decision on March 22, 2018; accepted on May 4, 2018

---

# Indexing k-mers

- Index all k-mers
  - Hash tables
    - Collision-free:  $O(4^k)$  possible k-mers
    - ... or many collisions
  - Other k-mer data structures (will follow)
- Minimizers:
  - “some” k-mers only
    - Reduced memory usage
  - Proposed in:
    - Roberts et al., “Reducing storage requirements for biological sequence comparison”, Bioinformatics, 2004

# Minimizers

- Ability to use sequence contiguity, while binning the k-mers
- Given a query length of  $N$ , all k-mers from the query will go to a different bin
- Minimizer-based binning fragments the query into  $\sim 3-4$  pieces
  - Representative k-mers
  - $(w,k)$ -minimizer: the smallest k-mer over  $w$  consecutive k-mers.

---

# Minimizers - example

Query (length 100):

```
ATGCGATATCGTAGGCGTCGATGGAGAGCTAGATCGATCGATCTAAATCCCGATCGATTCCGAGCGCGATCAAAGCGCGATA  
GGCTAGCTAAAGCTAGCA
```

Assume  $k=7$  - for the entire query (next slide)

---

# Minimizer, $k=7$ in the full query

| Position | 7-mer (forward strand) | 7-mer (reverse strand) |
|----------|------------------------|------------------------|
| 1        | ATGCGAT                | ATCGCAT                |
| 2        | TGCGATA                | TATCGCA                |
| 3        | GCGATAT                | ATATCGC                |
| 4        | CGATATC                | GATATCG                |
| 5        | GATATCG                | CGATATC                |
| 6        | ATATCGT                | ACGATAT                |
| 7        | TATCGTA                | TACGATA                |
| 8        | ATCGTAG                | CTACGAT                |
| 9        | TCGTAGG                | CCTACGA                |
| 10       | CGTAGGC                | GCCTACG                |
| 11       | GTAGGCG                | CGCCTAC                |
| 12       | TAGGCGT                | ACGCCTA                |
| .....    |                        |                        |
| 94       | GCTAGCA                | TGCTAGC                |

188 possible 7-mers in both strands -- smallest (alphabetical): 72nd k-mer in forward: **AAAGCGC**

# Minimizer, $k=7$ , $w=31$

| Position | 31-mer                                 | 7-mer minimizer |
|----------|--|-----------------|
| 1        | ATGCGATATCGTAGGCGTCGATGGAGAGCTAAGAGCTA | AGAGCTA         |
| 2        | TGCGATATCGTAGGCGTCGATGGAGAGCTAGAGAGCTA | AGAGCTA         |
| 3        | GCGATATCGTAGGCGTCGATGGAGAGCTAGAGAGCTA  | AGAGCTA         |
| 4        | CGATATCGTAGGCGTCGATGGAGAGCTAGATAGAGCTA | AGAGCTA         |
| 5        | GATATCGTAGGCGTCGATGGAGAGCTAGATCAGAGCTA | AGAGCTA         |
| 6        | ATATCGTAGGCGTCGATGGAGAGCTAGATCGAGAGCTA | AGAGCTA         |
| 7        | TATCGTAGGCGTCGATGGAGAGCTAGATCGAAGAGCTA | AGAGCTA         |
| 8        | ATCGTAGGCGTCGATGGAGAGCTAGATCGATAGAGCTA | AGAGCTA         |
| 9        | TCGTAGGCGTCGATGGAGAGCTAGATCGATCAGAGCTA | AGAGCTA         |
| 10       | CGTAGGCGTCGATGGAGAGCTAGATCGATCGAGAGCTA | AGAGCTA         |
| 11       | GTAGGCGTCGATGGAGAGCTAGATCGATCGAAGAGCTA | AGAGCTA         |
| 21       | ATGGAGAGCTAGATCGATCGATCTAAATCCC        | AAATCCC         |
| 22       | TGGAGAGCTAGATCGATCGATCTAAATCCCG        | AAATCCC         |
| 23       | GGAGAGCTAGATCGATCGATCTAAATCCCAG        | AAATCCC         |
| 24       | GAGAGCTAGATCGATCGATCTAAATCCCGAT        | AAATCCC         |
| 25       | AGAGCTAGATCGATCGATCTAAATCCCGATC        | AAATCCC         |
| 26       | GAGCTAGATCGATCGATCTAAATCCCGATCG        | AAATCCC         |
| 27       | AGCTAGATCGATCGATCTAAATCCCGATCGA        | AAATCCC         |
| 28       | GCTAGATCGATCGATCTAAATCCCGATCGAT        | AAATCCC         |

Minimizers= **AGAGCTA** and **AAATCCC**



# Minimizer-based binning

- A minimizer-based binning algorithm uses the minimizer of a k-mer to bin it. The algorithm works, because each k-mer has a unique minimizer of a certain size and will therefore always go to the same bin.
- K-mer based binning: bins for all possible kmers
- In the example in the previous slides,
  - K-mer based binning: 94 bins
  - Minimizer based binning: 5 bins

| Segment  | Minimizer-based bin |
|--|---------------------|
| ATGCGATATCGTAGGCGTCGATGGAGAGCTAGATCGATCGATCTAAATCC             | AGAGCTA             |
| ATGGAGAGCTAGATCGATCGATCTAAATCCCGATCGATTCCGAGCGCGATCAAAGAAATCCC |                     |
| AATCCCGATCGATTCCGAGCGCGATCAAAGC                                | AATCCCG             |
| ATCCCGATCGATTCCGAGCGCGATCAAAGCG                                | AATCGAT             |
| TCCCGATCGATTCCGAGCGCGATCAAAGCGCGATAGGCTAGCTAAAGCTAGCA          | AAAGCGC             |

# Minimizer-based mapping

- Minimap2
  - Collect minimizers from the reference sequence (*text*) and index in a hash table
  - For each query:
    - Collect minimizers from the query, find exact matches in the text to use as anchors
    - Identify co-linear anchors: chaining
    - Apply dynamic programming (DP) to extend from the ends of chains and to close regions between adjacent anchors in chains
      - Affine gap penalty

# Chaining

- An anchor is a 3-tuple  $(x, y, w)$ , indicating interval  $[x-w+1, x]$  on the reference matching interval  $[y-w+1, y]$  on the query.
- Given a list of anchors sorted by ending reference position  $x$ , let  $f(i)$  be the maximal chaining score up to the  $i^{\text{th}}$  anchor in the list.  $f(i)$  can be calculated with dynamic programming

$$f(i) = \max\left\{\max_{i > j \geq 1}\{f(j) + \alpha(j, i) - \beta(j, i)\}, w_i\right\}$$

# Chaining

$$f(i) = \max\{\max_{i>j\geq 1}\{f(j) + \alpha(j, i) - \beta(j, i)\}, w_i\}$$

$\alpha(j, i) = \min \{\min \{y_i - y_j, x_i - x_j\}, w_i\}$  : the number of matching bases between the two anchors.

$\beta(j, i) > 0$  : the gap cost



G: max dist between anchors

$$\beta(j, i) = \begin{cases} \infty, & \text{if } y_j \geq y_i. \\ \infty, & \max\{y_i - y_j, x_i - x_j\} > G. \\ \gamma_c((y_i - y_j) - (x_i - x_j)) & \text{otherwise} \end{cases}$$

$$\gamma_c(l) = \begin{cases} 0.01 \cdot \bar{w} \cdot |l| + 0.5 \log_2 |l| & (l \neq 0) \\ 0 & (l = 0) \end{cases}$$

w is the average seed length

# Chaining



---

# MAXIMAL UNIQUE MATCHES

---

# MUMs

- A *maximal unique match* of two strings  $S$  and  $T$  is a substring  $X$  that occurs exactly once in both  $S$  and  $T$  and cannot be extended to the left or to the right without losing one of the occurrences.

$$X = s_i s_{i+1} \dots s_{i+k-1}, X = t_j t_{j+1} \dots t_{j+k-1}, \text{ and } s_{i-1} \neq t_{j-1}, s_{i+k} \neq t_{j+k}$$

- To detect MUMs of two strings  $S$  and  $T$ , is simple
  - How?
  - And  $d$  strings  $(S^1, S^2, \dots, S^d)$ ?
    - in  $O(|\Sigma| + n)$  time where  $n = \sum_{i=1}^d |S^i|$

# MUMmer: whole genome alignment


- Whole mammalian genomes are ~3 Gb
  - Human chromosome 1 is ~250 Mb.
- Whole genome alignments using Smith-Waterman is therefore infeasible
- MUMmer (v1 1999, v2 2002, v3 2004) tries to solve this problem using MUMs as *anchors*
  - GPU accelerated version: MUMmerGPU (2013)
  - Combination of: suffix trees, longest increasing subsequences, and Smith-Waterman



# MUMmer


1. Find MUMs in Genomes A and B
2. Sort the matches found in the MUM alignment, and extract the longest possible set of matches that occur in the same order in both genomes.
  - a. LIS of a sequence of integers.
  - b. An ordered MUM alignment that provides an easy way to scan the alignment from left to right.
3. Close the gaps in the alignment by performing local identification of large inserts, repeats, small mutated regions, tandem repeats and substitutions
4. Output the alignment, including all the matches in the MUM alignment as well as the detailed alignments of regions that do not match exactly

# MUMmer




**Figure 3.** Aligning Genome *A* and Genome *B* after locating the MUMs. Each MUM is here indicated only by a number, regardless of its length. The top alignment shows all the MUMs. The shift of MUM 5 in Genome *B* indicates a transposition. The shift of MUM 3 could be simply a random match or part of an inexact repeat sequence. The bottom alignment shows just the LIS of MUMs in Genome *B*.

# MUMmer: *M. tuberculosis* strains




# MUMmer vs FASTA

FASTA



MUMmer



# MUMs for MSA

- If the sequences to align are similar, MUMs can be used as anchors for MSA
- Input:  $d$  sequences  $S^1, \dots, S^d$  of length  $n_1, \dots, n_d$
- Find max-length MUM for all sequences,  $\alpha$ 
  - This occurs exactly once in each  $S^i$  at position  $j_i$ , for  $1 \leq i \leq d$ , that you can use to split the set into two independent parts:

$$S^1_{1 \dots j_1}, S^2_{1 \dots j_2}, \dots, S^d_{1 \dots j_d} \text{ and } S^1_{j_1 + |\alpha| \dots n_1}, S^2_{j_2 + |\alpha| \dots n_2}, \dots, S^d_{j_d + |\alpha| \dots n_d}$$

- Apply the same recursively until sufficiently short sequences are obtained, suitable for MSA

---

# MAXIMAL EXACT MATCHES


---

# MEMs

- **MEMs**: exact matches between two strings that cannot be extended in either direction towards the beginning or end of two strings without allowing for a mismatch.
- Relaxation of MUMs, removing the uniqueness requirement
- The only constraint is the minimum length,  $L$
- All MEMs @ low memory → bidirectional BWT index

# BWA-MEM

- Seed-and-extend strategy for aligning billions of short queries to a large text (e.g., genome)
- Seeds -> Super-Maximal Exact Matches
  - SMEMs: MEMs that are not contained in other MEMs



*Vyverman et al.  
BMC  
Bioinformatics  
2015*



# BWA-MEM

1. Seeding: SMEMs: each query position the longest match covering the position
2. Chaining: greedy chaining algorithm, done while seeding
  - a. Discard short chains
3. Extension: banded affine-gap alignment